

Bilkent University  
Department of Computer Engineering

CS 342  
Operating Systems  
Spring 2003  
Final Exam

Date: May 28, 2003, Wednesday  
Time: 15:30  
Duration: 120 minutes

<b>Name</b>	
<b>Student ID</b>	

<b>Grade</b>	
--------------	--

**Please Read Carefully**

- Show your work and reasoning clearly!
- Write legibly!
- There should be total of 6 questions. Check your exam paper!
- You can use a calculator, but not a PDA or handheld computer.
- Answers that do not make sense will not get any credit!
- Try to write to the space provided under each question! You can use extra sheet(s) if that space is not enough. Do not forget to staple your extra sheet(s).

**Note that**

1 KB =  $2^{10}$  bytes  
1 MB =  $2^{20}$  bytes  
1 GB =  $2^{30}$  bytes

### **1. Problem (5 points)**

How much cylinder skew (in terms of number of sectors) is needed for a 7200 rpm disk with a track to track seek time of 2 msec? Assume the disk has 300 sectors in each track and each sector is 512 bytes.

X = The time that it takes until a whole sector passes under disk head =  $(1/7200) * 60 / 300 \text{ sec}$ .

Cylinder skew =  $0.002 \text{ sec} / (x \text{ sec/sector}) = \underline{\underline{72 \text{ sectors}}}$ .

## 2. Problem (16)

A machine that is using virtual memory has 32 bit virtual addresses and 256 MB physical memory. The page size is 4 KB. There are two programs, A and B, which will be run at different times. Program A has a size of 512MB and Program B has a size of 16 MB. Assume each page table entry requires 8 bytes of storage in memory.

- a) Compute the required page table size that each program (A and B) requires, if the system is using inverted page tables. Assume each program requires a different inverted page table and assume the whole physical memory is dedicated to the program (although this may not be realistic).

Memory size = 256 MB.

Number of frames in memory =  $256 \text{ MB} / 4 \text{ KB} = 65536 = 2^{16}$  page frames.

### 1) Program A

Inverted page table size = number of entries in the table \* entry size  
 $= 2^{16} * 2^3 = 2^{19} = 512 \text{ KB} = 524288 \text{ bytes} = 0.5 \text{ MB}$

### 2) Program B

The same with 1). 0.5 MB

- b) Compute the required page table size for each program if the system is using *two-level page tables*. Assume top-level page table index is 11 bits, and second-level page table index is 9 bits.

1) Program A

Each second level page table can address  $2^9 * 2^{12} = 2^{21} = 2 \text{ MB}$  address space.

For program A, we need  $512 \text{ MB} / 2 \text{ MB} = 256$  second level page table.

Total page table size =

Size of top level page table + 256 \* (size of a second level page table)

$$= (2^{11} * 2^3) + (2^8 * 2^9 * 2^3) = 2^{14} + 2^{20} = \underline{\underline{1064960 \text{ bytes}}}$$

2) Program B

Program B requires  $16 \text{ MB} / 2 \text{ MB} = 8$  second level page tables.

Total page table size =

Size of top level page table + 8 \* (size of a second level page table)

$$= (2^{11} * 2^3) + (2^3 * 2^9 * 2^3) = 2^{14} + 2^{15} = \underline{\underline{49152 \text{ bytes}}}$$

### 3. Problem (20)

Assume requests for disk blocks that are located in the **following cylinders** of a disk are queued: 55,58,39,18,90,160,150,38,184. Assume the requests are queued in the order of their arrivals. Compare the performance of the following disk scheduling algorithms using the tables (figures) provided in the next page. For each algorithm, at each step, compute the next cylinder visited, the number of cylinders that are traversed to reach to this next cylinder, and the total number of cylinders that are traversed (which is the total cost of the algorithm).

Make sure you write your final answers on the tables provided on the next page. Those tables will be graded!

- a) FIFO
- b) Shortest Seek First (SSF)
- b) Elevator (direction up) (starting in the direction of *increasing* cylinder number)
- c) Elevator (direction down) (starting in the direction of *decreasing* cylinder number)

Assume the head is initially at cylinder 100.

a) FIFO Starting at cylinder 100		
	Next cylinder accessed	Number of cylinders traversed
1	55	45
2	58	3
3	39	19
4	18	21
5	90	72
6	160	70
7	150	10
8	38	112
9	184	146
	Total cylinders traversed	498

b) SSF Starting at cylinder 100		
	Next cylinder accessed	Number of cylinders traversed
1	90	10
2	58	32
3	55	3
4	39	16
5	38	1
6	18	20
7	150	132
8	160	10
9	184	24
	Total cylinders traversed	248

c) Elevator (Up) Starting at cylinder 100		
	Next cylinder accessed	Number of cylinders traversed
1	150	50
2	160	10
3	184	24
4	90	94
5	58	32
6	55	3
7	39	16
8	38	1
9	18	20
	Total cylinders traversed	250

d) Elevator (Down) Starting at cylinder 100		
	Next cylinder accessed	Number of cylinders traversed
1	90	10
2	58	32
3	55	3
4	39	16
5	38	1
6	18	20
7	150	132
8	160	10
9	184	24
	Total cylinders traversed	248

#### 4. Problem (10 points)

Consider the two-dimensional array A:

```
int A[][] = new int[100][100];
```

where  $A[0][0]$  is stored at location 200, in a paged memory system with pages of size 200 bytes. A small process (whose size is 100 bytes) resides in page 0 (locations 0 through 99) for manipulating the matrix A; thus every instruction fetch will be from page 0.

Assume 3 page frames are allocated for the program (page frame 1, page frame 2, page frame 3). How many page faults will be generated by the following array-initialization loops, using LRU replacement, and assuming page frame 1 has the process in it, and the other two are initially empty.

**Information:** Assume a two dimensional array that is declared in a program is stored in virtual memory so that row 1 comes first, then row 2, then row 3, etc. In this case,  $A[0][0]$  is stored at virtual address 200,  $A[0][1]$  is stored at virtual address 204,  $A[1][0]$  is stored at virtual address 600, ,  $A[2][0]$  is stored at virtual address 1000, etc.

a)

```
for (int j = 0; j < 100; j++)  
    for (int i = 0; i < 100; i++)  
        A[i][j] = 0;
```

The reference string will look like the following (consists of the following page numbers).

:

First part:

1,3,5,...199, 1,3,5,...199,..... 1,3,5,...199

*(1,3,5,...199 string is repeated 50 times above. Each string contains 100 page numbers).*

Then, second part:

2,4,6,...,200, 2,4,6,...,200,..... ,2,4,6,...,200.

*(2,4,6,...200 string is repeated 50 times above. Each string contains 50 page numbers).*

Every new cell reference will cause a page fault in this case.

So first part will have:

100 x 50 = 5000 page faults.

Second part will have the same number of page faults:

Therefore, the total number of page faults will be : **10000** (ten thousand). .



b)

```
for (int i = 0; i < 100; i++)  
    for (int j = 0; j < 100; j++)  
        A[i][j] = 0;
```

Each page contains: 200 bytes / 4 = 50 integers.

Total number of pages required:  $100 \times 100 / 50 = 200$  pages.

A[0][0] is stored in page 1

A[0][1] is stored in page 1

A[0][2] is stored in page 1

...

A[0][49] is stored in page 1

A[0][50] is stored in page 2

A[0][51] is stored in page 2

....

A[0][99] is stored in page 2

A[1][0] is stored in page 3

...

A[0][150] is stored in page 4

...

A[0][1] is stored in page 4

Page accesses will be like the followin:

1,1,1,1...,1,2,2,2,...,2,3,3,3,...,3,.....199,199,199,...199,200,...2000

Each page is referenced 50 times. Since each page keeps 50 integers.

Since pages are accessed successively, there will be total of **200** page faults.

Every new page access will cause a page fault.

		FIFO Algorithm																											
Reference String <i>keeps the oldest page.</i>		1	0	2	2	1	7	6	7	0	1	2	0	3	0	4	5	1	5	2	4	5	6	7	6	7	2		
		1	1	1	1	1	1	0	0	0	2	2	7	6	6	1	0	3	3	4	4	4	5	1	1	1	1		
			0	0	0	0	0	2	2	2	7	7	6	1	1	0	3	4	4	5	5	5	1	2	2	2	2		
				2	2	2	2	7	7	7	6	6	1	0	0	3	4	5	5	1	1	1	2	6	6	6	6		
<i>keeps the newest page</i>							7	6	6	6	1	1	0	3	3	4	5	1	1	2	2	2	6	7	7	7	7		
Answer (Page Faults)		x	x	x			x	x			x		x	x		x	x	x		x			x	x					
Answer (number of page faults)		14																											

### b) LRU Algorithm

Reference String	LRU Algorithm																											
		1	0	2	2	1	7	6	7	0	1	2	0	3	0	4	5	1	5	2	4	5	6	7	6	7	2	
		1	0	2	2	1	7	6	7	0	1	2	0	3	0	4	5	1	5	2	4	5	6	7	6	7	2	
			1	0	0	2	1	7	6	7	0	1	2	0	3	0	4	5	1	5	2	4	5	6	7	6	7	
				1	1	0	2	1	1	6	7	0	1	2	2	3	0	4	4	1	5	2	4	5	5	5	6	
							0	2	2	1	6	7	7	1	1	2	3	0	0	4	1	1	2	4	4	4	5	
	Answer (PageFaults)		x	x	x			x	x		x		x		x		x	x	x		x			x	x			x

Answer (number of page faults)	15
--------------------------------	----

### c) NFU Algorithm

Reference String	NFU Algorithm																											
		1	0	2	2	1	7	6	7	0	1	2	0	3	0	4	5	1	5	2	4	5	6	7	6	7	2	
		1	1	2	2	2	2	2	7	7	1	2	2	2	0	0	0	1	1	2	2	2	2	2	2	7	2	
			0	1	1	1	1	1	2	2	7	1	1	1	2	2	2	0	0	1	1	1	1	1	1	2	7	
				0	0	0	7	7	1	1	2	7	0	0	1	1	1	2	5	0	0	0	0	0	0	1	1	
							0	6	6	0	0	0	7	3	3	4	5	5	2	5	4	5	6	7	6	0	0	
Answer (Page Faults)		x	x	x			x	x		x				x		x	x				x	x	x	x	x	x		

Answer (number of page faults)	15
--------------------------------	----

[illegible]

The above answer is based on the assumption that counters are never reset. I will also accept the answers that assume that counters will be reset when a page is evicted from memory. In this case, the answer will be:

	NFU Algorithm																											
Reference String		1	0	2	2	1	7	6	7	0	1	2	0	3	0	4	5	1	5	2	4	5	6	7	6	7	2	
Answer (Page Faults)		x	x	x			x	x		x				x	x	x	x				x	x	x					
Answer (number of page faults)												13																

## 6. Problem (25 points)

The following piece of code is given. Assume there is no buffering in I/O library and OS so that everything that is printed by `printf()` goes immediately to the screen without any delay. In the code, **`sleep(1)`** just causes the process to delay execution for 1 seconds. **`fflush(stdout)`** causes the output of `printf()` to be immediately flushed (written) to the screen without getting buffered.

```
#include <stdio.h>

void
main()
{
    int x, y;

    x = 20;
    y = 30;

    if (fork()==0)
    {
        y = 35;
        printf("%d\n", x); fflush(stdout);
        sleep(1);
        if (fork() == 0) {
            x = 25;
            printf("%d\n", x); fflush(stdout);
            printf("%d\n", y); fflush(stdout);
            sleep(1);
        }

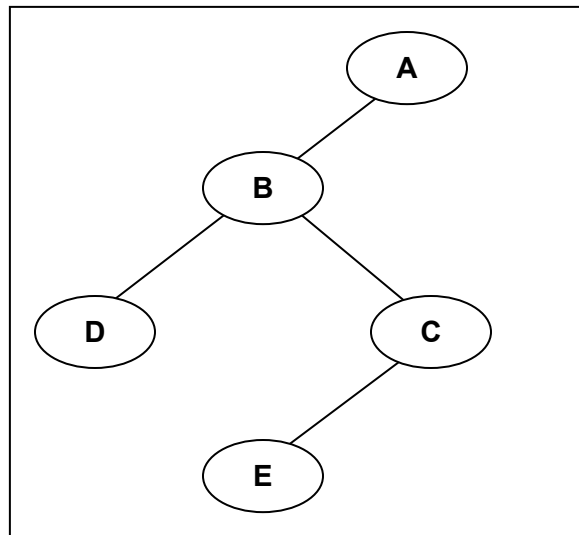
        if (fork() == 0)
        {
            printf("%d\n", x); fflush(stdout);
            exit(0);
            printf("%d\n", y); fflush(stdout);
        }
        else {
            printf("%d\n", x); fflush(stdout);
            printf("%d\n", y); fflush(stdout);
            sleep(1);
        }
    }
    printf("%d\n", x); fflush(stdout);
    printf("%d\n", y); fflush(stdout);
    exit(0);
}
```

Answer the following questions!. Please provide clean and clear answers!. Otherwise your answers will NOT get ANY credit.

- a) How many total number of processes will be created (including the first parent process that is created by typing the name of the program).

5 processes

- b) Draw the process tree (a tree that shows parent-child relationships). Give a letter to each process, such as A, B, C, D, E, F, G, .... A should be given to the initial parent process. You will use this identification scheme in some later questions.



Process Tree is shown above. The order of creation of processes are: A, B, C, D, E.

- c) The program prints numbers from the set {20, 25, 30, 35} to the screen. It prints one number at each line. How many numbers will be printed totally? (For example, if 20 20 25 would be printed to the screen – each number on a different line - , then the answer would be 3).

15

- d) How many 20's will be printed totally?

5

- e) How many 25's will be printed totally?

4

f) How many 30's will be printed totally?

1

g) How many 35's will be printed totally?

5

h) For each process that you have shown in the process tree (in question b) ), which numbers are printed by that process. If a number is printed many times, show that number many times (as much as it is printed).

A - 20 30

B - 20 20 35 20 35

C - 25 35 25 35 25 35

D - 20

E - 25

i) By referencing your answer in b) and using the identification of processes in that question, which process would terminate first.

A, since it does not execute a `sleep(1)` call. The scheduling granularity is in the order of 10s of milliseconds. During 1 second period, scheduler may intervene and schedule new processes lots of times.

j) By referencing your answer in b) and using the identification of processes in that question, which process would terminate last.

C, since it will call `sleep(1)` twice.